

Homework Assignment 2

Dillon Ford: A16092047

Imports

```
In [1]: from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
```

Images

```
In [2]: Figure_1 = Image.open('/Users/Dillon/Desktop/Winter_2020/COGS_118A/Assignment2/figure_1_decision_boundary_1.png')
Figure_2 = Image.open('/Users/Dillon/Desktop/Winter_2020/COGS_118A/Assignment2/figure_2_decision_boundary_2.png')
Figure_3 = Image.open('/Users/Dillon/Desktop/Winter_2020/COGS_118A/Assignment2/figure_3_decision_boundary_3.png')
Figure_4 = Image.open('/Users/Dillon/Desktop/Winter_2020/COGS_118A/Assignment2/figure_4_decision_boundary_4.png')
Figure_2_1 = Image.open('/Users/Dillon/Desktop/Winter_2020/COGS_118A/Assignment2/figure_2_1_my_decision_boundary.png')
Figure_2_4 = Image.open('/Users/Dillon/Desktop/Winter_2020/COGS_118A/Assignment2/figure_2_4_decision_boundary.png')
```

1. (10 points) Conceptual Questions

(1.1) Is the following statement true or false?

$f(x)$ is linear with respect to x , given $f(x) = w_0 + w_1x + w_2x^2$ where $x, w_0, w_1, w_2 \in R$.

False, a linear function, $f(x)$, is a polynomial function in which the variable x has a degree of at most one.

(1.2) “One-hot encoding” is a standard technique that turns categorical features into general real numbers. If we have a dataset S containing m data points where each data point has 1 categorical feature. Specifically, this categorical feature has k possible categories. Thus, the shape of the one-hot encoding matrix that represents the dataset S is:

- A. $k \times k$
- B. $1 \times k$
- C. $m \times k$ ←
- D. $m \times m$

Answer: C

(1.3) Assume we have a binary classification model:

$$f(x) = \begin{cases} +1, & w \cdot x + b \geq 0 \\ -1, & w \cdot x + b \leq 0 \end{cases}$$

where the feature vector $x = (x_1, x_2) \in \mathbb{R}^2$, bias $b \in \mathbb{R}$, weight vector $w = (w_1, w_2) \in \mathbb{R}^2$. The decision boundary of the classification model is:

$$w \cdot x + b = 0$$

(a) If the predictions of the classifier f and its decision boundary $w \cdot x + b = 0$ are shown in **Figure 1**, which one below can be a possible solution of weight vector w and bias b ?

In [3]: Figure_1

Out[3]:

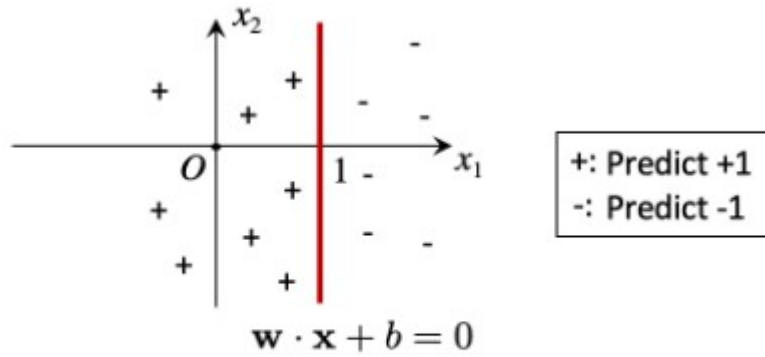


Figure 1: Decision Boundary 1

- A. $w = (+1, 0), b = -1$
- B. $w = (-1, 0), b = +1$ ←
- C. $w = (+1, 0), b = +1$
- D. $w = (0, -1), b = -1$

Answer: B

(b) If the predictions of the classifier f and its decision boundary $w \cdot x + b = 0$ are shown in **Figure 2**, which one below can be a possible solution of weight vector w and bias b ?

In [4]: Figure_2

Out[4]:

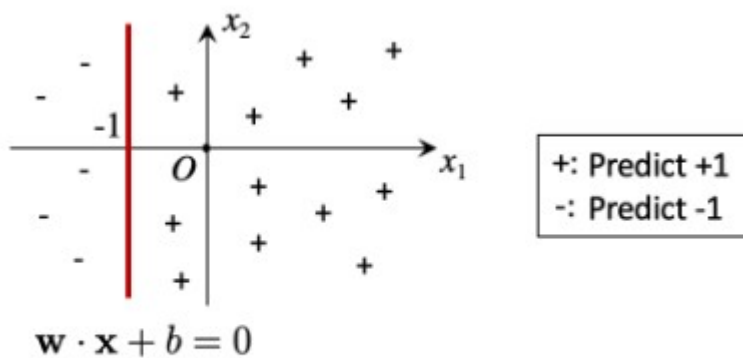


Figure 2: Decision Boundary 2

- A. $w = (+1, 0), b = -1$
- B. $w = (-1, 0), b = +1$
- C. $w = (+1, 0), b = +1$ ←
- D. $w = (0, -1), b = -1$

Answer: C

(1.4) Choose the most significant difference between regression and classification:

- A. unsupervised learning vs. supervised learning.
- B. $\boxed{\text{prediction of continuous values vs. prediction of class labels.}}$ ←
- C. features are not one-hot encoded vs features are one-hot encoded.
- D. none of the above.

Answer: B

2. (25 points) Decision Boundary

2.1 (3 points)

We are given a classifier that performs classification in R^2 (the space of data points with 2 features (x_1, x_2)) with the following decision rule:

$$h(x_1, x_2) = \begin{cases} 1, & \text{if } 2x_1 + 4x_2 - 8 \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

1. Draw the decision boundary of the classifier and shade the region where the classifier predicts 1. Make sure you have marked the x_1 and x_2 axes and the intercepts on those axes.

$$2x_1 + 4x_2 - 8 \geq 0$$

when $x_1 = 0$,

$$4x_2 - 8 = 0$$

$$\Rightarrow x_2 = 2$$

When $x_2 = 0$,

$$2x_1 - 8 = 0$$

$$\Rightarrow x_1 = 4$$

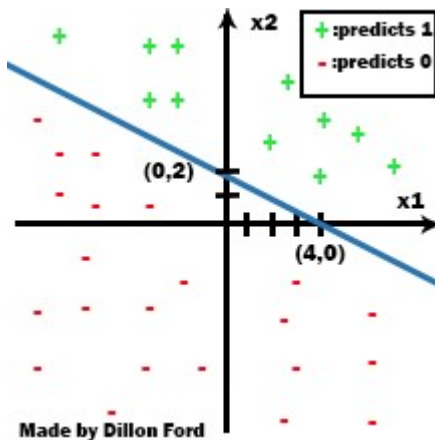
This gives us the intercepts, $(x_1, x_2) = (0, 2), (4, 0)$

As $2x_1 + 4x_2 - 8 \geq 0$ has the prediction 1
otherwise 0

We have the graph

In [5]: Figure_2_1

Out[5]:



2.2 (9 points)

We are given a classifier that performs classification in R^2 (the space of data points with 2 features (x_1, x_2)) with the following decision rule:

$$h(x_1, x_2) = \begin{cases} 1, & \text{if } w_1x_1 + w_2x_2 + b \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

Here, the normal vector w of the decision boundary is normalized, i.e.:

$$\|\mathbf{w}\|_2 = \sqrt{w_1^2 + w_2^2} = 1$$

1. Compute the parameters w_1 , w_2 and b for the decision boundary in **Figure 3**. Please make sure the predictions from the obtained classifier are consistent with **Figure 3**.

Hint: Please use the intercepts in the **Figure 3** to find the relation between w_1 , w_2 and b . Then, substitute it into the normalization constraint to solve for parameters.

```
In [6]: Figure_3
```

```
Out[6]:
```

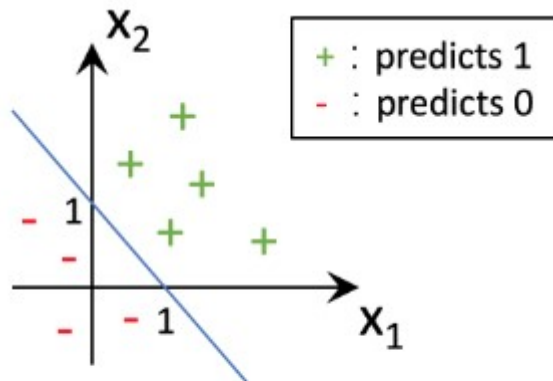


Figure 3: Decision boundary to solve for parameters.

From the graph we can see two intercepts, one on the x_1 axis and one on the x_2 axis

Where $x_1 = 1$ and $x_2 = 0$ (point on x_1 axis) i.e. the point $(1, 0)$

Where $x_1 = 0$ and $x_2 = 1$ (point on x_1 axis) i.e. the point $(0, 1)$

For, $w_1 x_1 + w_2 x_2 + b = 0$

Substituting, $(1, 0) \Rightarrow w_1 + b = 0 \dots(1)$

Substituting, $(0, 1) \Rightarrow w_2 + b = 0 \dots(2)$

Subtracting equations (1) and (2) yields,

$$w_1 = w_2$$

Since, $\sqrt{w_1^2 + w_2^2} = 1$, and $w_1 = w_2$ then it becomes,

$$\sqrt{2w_1^2} = 1$$

$$\Rightarrow 2w_1^2 = 1$$

$$\Rightarrow w_1 = w_2 = \frac{1}{\sqrt{2}}$$

Solving for b from equation (1) or (2) yields,

$$\frac{1}{\sqrt{2}} + b = 0$$

$$\Rightarrow b = -\frac{1}{\sqrt{2}}$$

Hence our parameters are,

$$\boxed{w_1 = w_2 = \frac{1}{\sqrt{2}} \text{ and } b = -\frac{1}{\sqrt{2}}} \leftarrow$$

2. Use parameters from the above question to compute predictions for the following two data points:

$A = (3, 6), B = (1, -4)$.

For the data point $A = (3, 6)$ we have,

$$\frac{3}{\sqrt{2}} + \frac{6}{\sqrt{2}} - \frac{1}{\sqrt{2}} \geq 0$$

$$\Rightarrow 4\sqrt{2} \geq 0$$

The predicted label for the data point $A = 1$ ←

For the data point $B = (1, -4)$ we have,

$$\frac{1}{\sqrt{2}} - \frac{4}{\sqrt{2}} - \frac{1}{\sqrt{2}} \geq 0$$

$$\Rightarrow -2\sqrt{2} \not\geq 0$$

The predicted label for the data point $B = 0$ ←

2.3 (10 points)

We are given a classifier that performs classification in R^3 (the space of data points with 3 features (x_1, x_2, x_3)) with the following decision rule:

$$h(x_1, x_2, x_3) = \begin{cases} 1, & \text{if } w_1x_1 + w_2x_2 + w_3x_3 + b \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

Here, the normal vector w of the decision boundary is normalized, i.e.:

$$\|w\|_2 = \sqrt{w_1^2 + w_2^2 + w_3^2} = 1$$

In addition, we set $b \leq 0$ to have a unique equation for the decision boundary.

1. Compute the parameters w_1, w_2, w_3 and b for the decision boundary that passes through three points $A = (3, 2, 4), B = (-1, 0, 2), C = (4, 1, 5)$ in **Figure 4**.

Hint: Please use the intercepts in the **Figure 4** to find the relation between w_1, w_2, w_3 and b . Then, substitute it into the normalization constraint to solve for parameters.


```
In [7]: Figure_4
```

```
Out[7]:
```

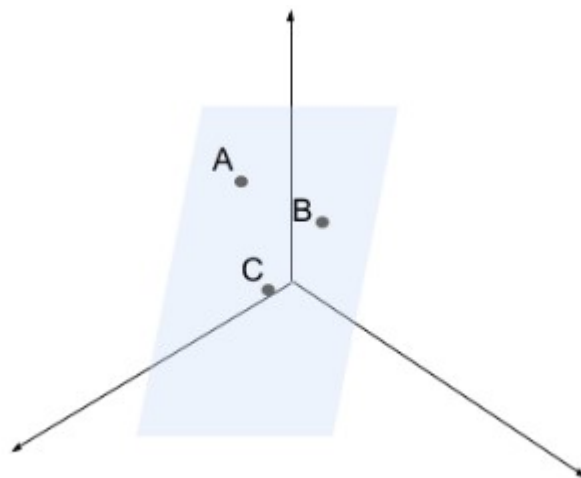


Figure 4: Decision boundary to solve the parameters.

Equation of a plane passing through,

$$A = (3, 2, 4)$$

$$B = (-1, 0, 2)$$

$$C = (4, 1, 5)$$

$$B - A = (-4, -2, -2) = \vec{v}_1$$

$$C - B = (5, 1, 3) = \vec{v}_2$$

$$\vec{v}_1 \times \vec{v}_2 = (-2 \cdot 3 - (-2 \cdot 1) - 2 \cdot 5 - 4 \cdot 1 - (-2 \cdot 5)) = (-4, 2, 6)$$

$$\Rightarrow (w_1, w_2, w_3) = \frac{(-4, 2, 6)}{\sqrt{4^2 + 2^2 + 6^2}} = \frac{-4}{\sqrt{56}}, \frac{2}{\sqrt{56}}, \frac{6}{\sqrt{56}}$$

$$\sqrt{\frac{-4}{\sqrt{56}}^2 + \frac{2}{\sqrt{56}}^2 + \frac{6}{\sqrt{56}}^2} = 1$$

$$\Rightarrow \sqrt{\frac{2}{7} + \frac{1}{14} + \frac{9}{14}} = 1$$

$$\Rightarrow \sqrt{1} = 1 \Rightarrow 1 = 1$$

$$\Rightarrow w_1 = \frac{-4}{\sqrt{56}}, w_2 = \frac{2}{\sqrt{56}}, w_3 = \frac{6}{\sqrt{56}} \leftarrow$$

$$\text{as, } B = (-1, 0, 2) \Rightarrow (x_1, x_2, x_3)$$

$$\Rightarrow w_1 x_1 + w_2 x_2 + w_3 x_3 = -b$$

$$\Rightarrow \frac{4}{\sqrt{56}} + 0 + \frac{12}{\sqrt{56}} = -b$$

$$\Rightarrow b = \frac{-16}{\sqrt{56}} \leftarrow$$

2. Use parameters from the above question to compute predictions for the following two data points:

$$D = (0, 0, 0), E = (1, 0, 5).$$

For the data point $D = (0, 0, 0)$ we have,

$$0 + 0 + 0 - \frac{-16}{\sqrt{56}}$$

$$\Rightarrow -\frac{-16}{\sqrt{56}} \not\geq 0$$

The predicted label for the data point $D = 0$ ←

For the data point $E = (1, 0, 5)$ we have,

$$\frac{-4}{\sqrt{56}} + 0 + \frac{30}{\sqrt{56}} - \frac{16}{\sqrt{56}}$$

$$\Rightarrow \frac{10}{\sqrt{56}} \geq 0$$

The predicted label for the data point $E = 1$ ←

2.4 (3 points)

We are given a classifier that performs classification in R^2 (the space of data points with 2 features (x_1, x_2)) with the following decision rule:

$$h(x_1, x_2) = \begin{cases} 1, & \text{if } x_1^2 + x_2^2 - 10 \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

1. Draw the decision boundary of the classifier and shade the region where the classifier predicts 1. Make sure you have marked the x_1 and x_2 axes and the intercepts on those axes.

$$x_1^2 + x_2^2 - 10 \geq 0$$

when $x_1 = 0$,

$$\Rightarrow x_2 = \sqrt{10}$$

When $x_2 = 0$,

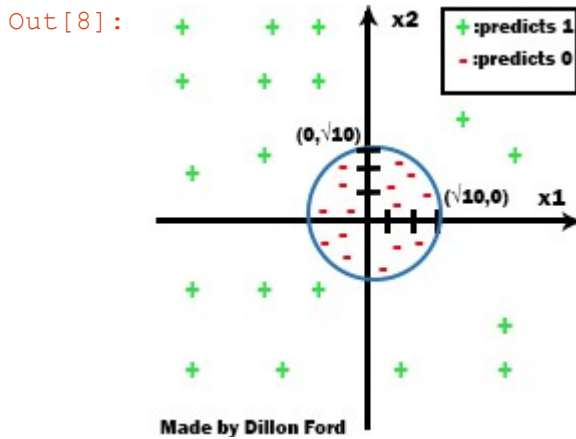
$$\Rightarrow x_1 = \sqrt{10}$$

This gives us the intercepts, $(x_1, x_2) = (0, \sqrt{10}), (\sqrt{10}, 0)$

As $x_1^2 + x_2^2 - 10 \geq 0$ has the prediction 1
otherwise 0

We have the graph

In [8]: Figure_2_4



3. (10 points) Derivatives

3.1 Function Defined by Scalars

- (3 points) Given a function $f(w) = (y_1 + wx_1)^2$ where $(x_1, y_1) = (3, 4)$ represents a data point, derive $\frac{\partial f(w)}{\partial w}$

$$\frac{\partial f(w)}{\partial w} = 2(y_1 + wx_1)(x_1)$$

such that, $(x_1, y_1) = (3, 4)$ we have,

$$\frac{\partial f(w)}{\partial w} = 2(4 + 3w)(3)$$

$$\boxed{\frac{\partial f(w)}{\partial w} = 6(4 + 3w)} \leftarrow$$

2. (3 points) Given a function $f(w) = \sum_{i \in [1,2]} (y_i - wx_i)^2$ where $(x_1, y_1) = (1, 1)$, $(x_2, y_2) = (2, 3)$ are two data points, derive $\frac{\partial f(w)}{\partial w}$

$$f(w) = (y_1 - wx_1)^2 + (y_2 - wx_2)^2$$

$$\frac{\partial f(w)}{\partial w} = 2(y_1 - wx_1)(-x_1) + 2(y_2 - wx_2)(-x_2)$$

For $(x_1, y_1) = (1, 1)$ and $(x_2, y_2) = (2, 3)$ we have,

$$\frac{\partial f(w)}{\partial w} = 2(1 - w)(-1) + 2(3 - 2w)(-2)$$

$$\boxed{\frac{\partial f(w)}{\partial w} = 2(w - 1) + 4(2w - 3)} \leftarrow$$

3.2 Function Defined by Vectors

1. (4 points) Given a function $f(w) = (y - wx)^T(y - wx)$ where $x = [1, 2]^T$ and $y = [1, 3]^T$, derive $\frac{\partial f(w)}{\partial w}$.

Note: In $f(w)$, $w \in \mathbb{R}$ is still a scalar.

$$f(w) = \left(\begin{bmatrix} 1 \\ 3 \end{bmatrix} - w \begin{bmatrix} 1 \\ 2 \end{bmatrix} \right)^T \left(\begin{bmatrix} 1 \\ 3 \end{bmatrix} - w \begin{bmatrix} 1 \\ 2 \end{bmatrix} \right)$$

$$f(w) = \left(\begin{bmatrix} 1 \\ 3 \end{bmatrix} - \begin{bmatrix} w \\ 2w \end{bmatrix} \right)^T \left(\begin{bmatrix} 1 \\ 3 \end{bmatrix} - \begin{bmatrix} w \\ 2w \end{bmatrix} \right)$$

$$f(w) = \begin{bmatrix} 1 - w \\ 3 - 2w \end{bmatrix}^T \begin{bmatrix} 1 - w \\ 3 - 2w \end{bmatrix}$$

$$f(w) = [1 - w, 3 - 2w] \begin{bmatrix} 1 - w \\ 3 - 2w \end{bmatrix}$$

$$f(w) = (1 - w)(1 - w) + (3 - 2w)(3 - 2w)$$

$$f(w) = (1 - w)^2 + (3 - 2w)^2$$

$$f(w) = (1 - 2w + w^2) + (9 - 12w + 4w^2)$$

$$f(w) = 10 - 14w + 5w^2$$

$$\frac{\partial f(w)}{\partial w} = 0 - 14 + 10w$$

$$\boxed{\frac{\partial f(w)}{\partial w} = 10w - 14} \leftarrow$$

4. (9 points) Concepts

Select the correct option(s). Note that there might be multiple correct options

1. For two monotonically increasing functions $f(x)$ and $g(x)$:

A. $f(x) + g(x)$ is always monotonically increasing. \leftarrow

B. $f(x) - g(x)$ is always monotonically increasing.

C. $f(x^2)$ is always monotonically increasing.

D. $f(x^3)$ is always monotonically increasing. \leftarrow

Answer: A, D

2. For a function $f(x) = x(10 - x)$, $x \in \mathbb{R}$, please choose the correct statement(s) below:

$$f'(x) = 10 - 2x = 0$$

$$\Rightarrow x = 5, > 0 \Rightarrow \text{maxima at } x = 5$$

$$f''(x) = -2, < 0 \Rightarrow \text{no minima}$$

$$\operatorname{argmax}_x f(x) = 5$$

$$\max_x f(x) = f(5) = 5 \cdot (5) = 25$$

A. $\boxed{\operatorname{argmax}_x f(x) = 5.} \leftarrow$

B. $\operatorname{argmin}_x f(x) = 25.$

C. $\min_x f(x) = 5.$

D. $\boxed{\max_x f(x) = 25.} \leftarrow$

Answer: A,D

3. Assume we have a function $f(x)$ which is differentiable at every $x \in \mathbb{R}$. There are three properties that describe the function $f(x)$:

(1) $f(x)$ is a convex function.

(2) When $x = x_0$, $f'(x_0) = 0$.

(3) $f(x_0)$ is a global minimum of $f(x)$.

Which one of the following statements is wrong?

Hint: You can use a failure case to disprove a statement.

A. Given (1) and (2), we can prove that (3) holds.

B. $\boxed{\text{Given (2) and (3), we can prove that (1) holds.}} \leftarrow$

C. Given (1) and (3), we can prove that (2) holds.

Answer: B

5 (4 points) Argmin and Argmax

An unknown estimator is given an estimation problem to find the minimizer and maximizer of the objective function $G(w) \in (0, 2]$:

$$(w_a, w_b) = (\operatorname{argmin}_w G(w), \operatorname{argmax}_w G(w)) \dots (1)$$

The solution to Eq.1 by the estimator is $(w_a, w_b) = (10, 20)$.

Given this information, please obtain the value of w^* such that:

$$w^* = \operatorname{argmin}_w [10 - 4 \times \ln(G(w))] \dots (2)$$

$\ln(x)$ is a monotonic increasing function with $\ln(x) > 0$ for $x > 0$

$\ln(G(w))$ is a monotonic increasing function

$-\ln(x)$ is a monotonic decreasing function

$-4\ln(x)$ is a monotonic decreasing function

so, $10 - 4 \times \ln(x)$ is minimum when $G(w)$ is maximum.

$(w_a, w_b) = (10, 20)$ and so,

$$\boxed{w^* = w_b = 20} \leftarrow$$

6. (12 points) Data Manipulation

In this question, we still use the Iris dataset from Homework 1. In fact, you can see the shape of array X is (150,4) by running **X.shape**, which means it contains 150 data points whereeach has 4 features. Here, we will perform some basic data manipulation and calculate some statistics:

1. Divide array X evenly to five subsets of data points:

Group 1: 1st to 30th data point,

Group 2: 31st to 60th data point,

Group 3: 61st to 90th data point,

Group 4: 91st to 120th data point,

Group 5: 121st to 150th data point.

Then calculate the mean of feature vectors in each group. Your results should be five 4-dimensional vectors (i.e. shape of NumPy array can be (4, 1), (1, 4) or (4,))

```
In [9]: iris = datasets.load_iris()
        X = iris.data
        Y = iris.target
```



```
In [10]: # divide the dataset into five subsets
group_1 = X[0:30,:]
group_2 = X[30:60,:]
group_3 = X[60:90,:]
group_4 = X[90:120,:]
group_5 = X[120:150,:]
print('Shape of group subsets:',group_1.shape, group_2.shape, group_3.s
hape, group_4.shape, group_5.shape)
```

Shape of group subsets: (30, 4) (30, 4) (30, 4) (30, 4) (30, 4)

```
In [11]: # calculate the mean of feature vectors in each group
group_1_mean = np.mean(group_1, axis=0)
group_2_mean = np.mean(group_2, axis=0)
group_3_mean = np.mean(group_3, axis=0)
group_4_mean = np.mean(group_4, axis=0)
group_5_mean = np.mean(group_5, axis=0)
print('Mean of Group 1 feature vectors:',group_1_mean )
print('Mean of Group 2 feature vectors:',group_2_mean )
print('Mean of Group 3 feature vectors:',group_3_mean )
print('Mean of Group 4 feature vectors:',group_4_mean )
print('Mean of Group 5 feature vectors:',group_5_mean )
```

Mean of Group 1 feature vectors: [5.02666667 3.45 1.47333333 0.24666667]
Mean of Group 2 feature vectors: [5.35 3.22 2.42 0.62333333]
Mean of Group 3 feature vectors: [5.98 2.75 4.3 1.34]
Mean of Group 4 feature vectors: [6.25333333 2.85666667 5.11333333 1.77333333]
Mean of Group 5 feature vectors: [6.60666667 3.01 5.48333333 2.01333333]

2. Remove 2nd and 3rd features from array X, resulting a 150×2 matrix. Then calculate the mean of all feature vectors. Your result should be a 2-dimensional vector.

```
In [12]: X_1 = X[:, [0,3]]
print('shape of X:',X_1.shape)
print('mean of feature vectors:',np.mean(X_1, axis=0))
```

shape of X: (150, 2)
mean of feature vectors: [5.84333333 1.19933333]

3. Remove last 10 data points from array X, resulting a 140×4 matrix. Then calculate the mean of feature vectors. Your result should be a 4-dimensional vector.

```
In [13]: X_2 = X[:-10,:]
print('shape of X:', X_2.shape)
print('mean of feature vectors:', np.mean(X_2, axis=0))

shape of X: (140, 4)
mean of feature vectors: [5.8          3.05928571  3.64571429  1.13
]
```

7. (15 points) Training vs. Testing Errors

In this problem, we are given two trained predictive models on a modified Iris dataset. Each data point (x, y) has a feature vector $x_i \in \mathbb{R}^4$ and its corresponding label $y_i \in [0, 1]$, where $i \in [1, 2, \dots, 150]$. To predict on the new data, here we consider two types of model: a regression model and a classification model. The regression model is trained to predict a real number, while the classification model applies a threshold to the output of the regression model, converting the real number into a binary value.

The regression model is as followed:

$$\hat{y}_i(x_i) = w^T x_i + b$$

The classifier is as followed:

$$h(x_i) = \begin{cases} 1, & \text{if } \hat{y}_i(x_i) \geq \frac{1}{2} \\ 0, & \text{otherwise} \end{cases}$$

where $w = [0.1297, 0.1225, -0.1171, 0.6710]^T$, $b = -1.1699$.

The regression error is defined as:

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2}$$

and the classification error is defined as:

$$\frac{1}{n} \sum_{i=1}^n 1(h(x_i) \neq y_i)$$

where n is the number of data points.

The data as well as the split of training and testing set are given in the Jupyter notebook we provided. **You should not use the scikit-learn library.**

Please download the notebook `training_test_errors.ipynb` from the course website and fill in the missing blanks. Follow the instructions in the skeleton code and report:

- Training error of the regression model.
- Testing error of the regression model.
- Training error of the classification model.
- Testing error of the classification model.

Load the Iris dataset

```
In [14]: # Iris dataset.
iris = datasets.load_iris()      # Load Iris dataset.

X = iris.data                    # The shape of X is (150, 4), which means
                                # there are 150 data points, each data
                                # point
                                # has 4 features.

# Here for convenience, we divide the 3 kinds of flowers into 2 groups:
#     Y = 0 (or False): Setosa (original value 0) / Versicolor (original value 1)
#     Y = 1 (or True):  Virginica (original value 2)

# Thus we use (iris.target > 1.5) to divide the targets into 2 groups.
# This line of code will assign:
#     Y[i] = True  (which is equivalent to 1) if iris.target[k] > 1.5 (Virginica)
#     Y[i] = False (which is equivalent to 0) if iris.target[k] <= 1.5 (Setosa / Versicolor)

Y = (iris.target > 1.5).reshape(-1,1) # The shape of Y is (150, 1), which means
                                      # there are 150 data points, each data
                                      # point
                                      # has 1 target value.

X_and_Y = np.hstack((X, Y))        # Stack them together for shuffling.
np.random.seed(1)                  # Set the random seed.
np.random.shuffle(X_and_Y)         # Shuffle the data points in X_and_Y array

print(X.shape)
print(Y.shape)
print(X_and_Y[0])                  # The result should be always: [ 5.8
4.   1.2  0.2  0. ]

(150, 4)
(150, 1)
[5.8 4.   1.2 0.2 0. ]
```

In [15]: *# Divide the data points into training set and test set.*

```
X_shuffled = X_and_Y[:, :4]
```

```
Y_shuffled = X_and_Y[:, 4]
```

```
X_train = X_shuffled[:100] # Shape: (100,4)
```

```
Y_train = Y_shuffled[:100] # Shape: (100,)
```

```
X_test = X_shuffled[100:] # Shape: (50,4)
```

```
Y_test = Y_shuffled[100:] # Shape: (50,)
```

```
print(X_train.shape)
```

```
print(Y_train.shape)
```

```
print(X_test.shape)
```

```
print(Y_test.shape)
```

```
(100, 4)
```

```
(100,)
```

```
(50, 4)
```

```
(50,)
```

In [16]: **from** sklearn.linear_model **import** LinearRegression

```
# let's train a LR model...
```

```
# note that this time we let sklearn fit the intercept
```

```
# ask yourself why. Ask yourself what could we have done
```

```
# to X_and_Y so that we could have used
```

```
# = LinearRegression(fit_intercept=False).fit(X_train, Y_train)
```

```
# instead of the below line, and got identical results?
```

```
pre_defined_weights = LinearRegression().fit(X_train, Y_train)
```

```
w = pre_defined_weights.coef_
```

```
b = pre_defined_weights.intercept_
```

```
In [20]: def regression_error(x, y, w, b):

    regression_error = 0
    for i in range(len(x)):

        # TODO: ***** To be filled *****

        # prediction based on x

        y_hat = np.dot(w.transpose(), x[i]) + b

        # regression error, doing the sum
        regression_error += np.square(y_hat-y[i])

    # calculate the mean and square root
    regression_error = np.sqrt(regression_error*(1/len(x)))

    return regression_error

def classification_error(x, y, w, b):
    classification_error = 0

    for i in range(len(x)):

        # TODO: ***** To be filled *****

        # prediction based on x
        h_x = 0

        y_hat = np.dot(w.transpose(), x[i]) + b

        # classification error
        if y_hat >= (1/2):
            h_x = 1
        error = 1*int(h_x != y[i])

        classification_error += error

    # calculate the mean of error
    classification_error = (classification_error)/(len(x))

    return classification_error

print('Training regression errors are:')
print(regression_error(X_train, Y_train, w, b))
print('Testing regression errors are:')
print(regression_error(X_test, Y_test, w, b))

print('Training classification errors are:')
print(classification_error(X_train, Y_train, w, b))
print('Testing classification errors are:')
print(classification_error(X_test, Y_test, w, b))
```

```

Training regression errors are:
0.2792069270624264
Testing regression errors are:
0.33046623492235216
Training classification errors are:
0.06
Testing classification errors are:
0.14

```

8. (15 points) Linear Regression

Assume we are given a dataset $S = \{(x_i, y_i), i = 1, \dots, n\}$. Here, $x_i \in \mathcal{R}$ is a feature scalar (a.k.a. value of input variable) and $y_i \in \mathcal{R}$ is its corresponding value (a.k.a. value of dependent variable). In this section, we aim to fit data points with a line:

$$y = w_0 + w_1 x \quad (3)$$

where $w_0, w_1 \in \mathcal{R}$ are two parameters to determine the line. Next, we measure the quality of fitting by evaluating a sum-of-squares error function $g(w_0, w_1)$:

$$g(w_0, w_1) = \sum_{i=1}^n (w_0 + w_1 x_i - y_i)^2 \quad (4)$$

When $g(w_0, w_1)$ is near zero, it means the proposed line can fit the dataset and model an accurate relation between x_i and y_i . The best line with parameters (w_0^*, w_1^*) can reach the minimum value of the error function $g(w_0, w_1)$:

$$(w_0^*, w_1^*) = \operatorname{argmin} g(w_0, w_1) \quad (5)$$

To obtain the parameters of the best line, we will take the gradient of function $g(w_0, w_1)$ and set it to zero. That is:

$$\nabla g(w_0, w_1) = 0 \quad (6)$$

The solution (w_0^*, w_1^*) of the above equation will determine the best line $y = w_0^* + w_1^* x$ that fits the dataset S .

In reality, we typically tackle this task in a matrix form: First, we represent data points as matrices

$X = [x_1, x_2, \dots, x_n]^T$ and $Y = [y_1, y_2, \dots, y_n]^T$, where $x_i = [1, x_i]^T$ is a feature vector corresponding to x_i . The parameters of the line are also represented as a matrix $W = [w_0, w_1]^T$. Thus, the sum-of-squares error function $g(W)$ can be defined as (a.k.a. squared L_2 norm):

$$g(W) = \sum_{i=1}^n (x_i^T W - y_i)^2 \quad (7)$$

$$= \|XW - Y\|_2^2 \quad (8)$$

$$= (XW - Y)^T (XW - Y) \quad (9)$$

Similarly, the parameters $W^* = [w_0^*, w_1^*]^T$ of the best line can be obtained by solving the equation below:

$$\nabla g(W) = \frac{\partial g(W)}{\partial W} = 0 \quad (10)$$

(a) According to Eq. 8 and 9, compute the gradient of $g(W)$ with respect to W . Your result should be in the form of X , Y and W .

$$\begin{aligned}
g(W) &= (XW - Y)^T(XW - Y) \\
g(W) &= (X^T W^T - Y^T)(XW - Y) \\
g(W) &= X^T W^T XW - W^T X^T Y - Y^T XW + Y^T Y
\end{aligned}$$

Since, $W^T X^T Y$ and $Y^T XW$ are 1x1 matrices,

$$\begin{aligned}
(W^T X^T Y)^T &= Y^T XW \\
\Rightarrow W^T X^T Y &= Y^T XW
\end{aligned}$$

$$\text{Thus, } g(W) = Y^T Y - 2W^T X^T Y + W^T X^T XW$$

$$\begin{aligned}
\nabla g(W) &= \nabla(Y^T Y - 2W^T X^T Y + W^T X^T XW) \\
\nabla g(W) &= \nabla(Y^T Y) - 2\nabla(W^T X^T Y) + \nabla(W^T X^T XW)
\end{aligned}$$

$$\begin{aligned}
\nabla(Y^T Y) &\Rightarrow \frac{\partial}{\partial W}(Y^T Y) = 0 \\
\nabla(W^T X^T Y) &\Rightarrow \frac{\partial}{\partial W}(W^T X^T Y) = X^T Y \\
\nabla(W^T X^T XW) &\Rightarrow \frac{\partial}{\partial W}(W^T X^T XW) = 2X^T XW
\end{aligned}$$

$$\begin{aligned}
\nabla g(W) &= 0 - 2X^T Y + 2X^T XW \\
\nabla g(W) &= 2(X^T XW - X^T Y) \leftarrow
\end{aligned}$$

(b) By setting the answer of part (a) to 0, prove the following:

$$W^* = \operatorname{argmin}_w g(W) = (X^T X)^{-1} X^T Y \quad (11)$$

Note: The above formula demonstrates a closed form solution of Eq. 10

$$\begin{aligned}
\nabla g(W) &= 0 \\
\Rightarrow 2(X^T XW - X^T Y) &= 0 \\
\Rightarrow X^T XW - X^T Y &= 0 \\
\Rightarrow X^T XW &= X^T Y \\
\Rightarrow (X^T X)^{-1} X^T XW &= (X^T X)^{-1} X^T Y \\
\Rightarrow W &= (X^T X)^{-1} X^T Y
\end{aligned}$$

So, the parameter of our best line is given by,

$$W^* = (X^T X)^{-1} X^T Y \leftarrow$$

Previously, we define a sum-of-squares error function $g(w_0, w_1) = \sum_{i=1}^n (y_i - w_0 - w_1 x_i)^2$ and represent it in a matrix form $g(W) = \|XW - Y\|_2^2$. Actually, we can have multiple choices of the error function: For example, we can define a sum-of-absolute error function $h(w_0, w_1)$:

$$h(w_0, w_1) = \sum_{i=1}^n |w_0 + w_1 x_i - y_i| \quad (12)$$

and represent it in a matrix form $h(W)$ (a.k.a. L_1 norm):

$$h(W) = \sum_{i=1}^n |x_i^T W - y_i| \quad (13)$$

$$= \|XW - Y\|_1 \quad (14)$$

(c) According to the Eq. 13, compute the gradient of the error function $h(W)$ with respect to W . Your result should be in the form of x_i, y_i and W .

Hint: Given a function $f(x) \in \mathbb{R}$, we have:

$$\frac{\partial |f(x)|}{\partial W} = \text{sign}(f(x)) \frac{\partial f(x)}{\partial x}$$

where,

$$\text{sign}(x) = \begin{cases} 1, x > 0 \\ 0, x = 0 \\ -1, x < 0 \end{cases}$$

$$h(W) = \sum_{i=1}^n |x_i^T W - y_i|$$

$$\nabla h(W) = \frac{\partial |h(W)|}{\partial W} = \sum_{i=1}^n \text{sign}(x_i^T W - y_i) \frac{\partial h(W)}{\partial W}$$

$$\boxed{\nabla h(W) = \sum_{i=1}^n (\text{sign}(x_i^T W - y_i)) x_i} \leftarrow$$