

# Homework Assignment 6

COGS 118A: Supervised Machine Learning Algorithms

**Due: Dec 6th, 2020, 11:59pm (Pacific Time).**

**Instructions:** Combine your answer to the questions below with your notebook to create a PDF file; submit your file via Gradescope.

## 1 (8 points) Conceptual Questions

1. Which one below best describes what support vectors are:
  - A. The decision boundary.
  - B. The positive and the negative planes.
  - C. The training samples that determine the positive and the negative planes.
  - D. The test samples that determine the positive and the negative planes.
  
2. When tuning the hyper-parameters of a model, we find the hyper-parameters that
  - A. minimize error on the test set
  - B. minimize error on the test set
  - C. maximize the margin of the classifier
  - D. minimize error on the validation set
  
3. Select **all** that apply. k-fold cross validation is
  - A. Not necessary when you have huge amounts of labelled data
  - B. A way to do model selection while minimizing over-fitting

- C. The only way to do hyper-parameter tuning
  - D. Subject to the bias-variance trade-off
- 
- 4. When you increase the  $k$  in cross-validation while keeping the dataset the same, you
    - A. Get a decrease in both the generalization error and the validation error
    - B. Get an increase both the generalization error and the validation error
    - C. Get a decrease in the variability of the validation error across folds
    - D. Get an increase in the variability of the validation error across folds

## 2 (10 points) Cross-Validation

Given a training dataset  $S_{\text{training}} = \{(x_i, y_i)\}, i = 1, \dots, 6\}$  where  $x_i \in \mathbb{R}$  is the feature scalar and  $y_i \in \{-1, +1\}$  is the corresponding label. The data points in the dataset  $S_{\text{training}}$  are given below:

$$(x_1, y_1) = (2, -1), \quad (x_2, y_2) = (7, -1), \quad (x_3, y_3) = (4, +1),$$

$$(x_4, y_4) = (1, -1), \quad (x_5, y_5) = (3, +1), \quad (x_6, y_6) = (6, +1).$$

Suppose you are training a linear classifier  $f(x; a, b) = \text{sign}(ax + b)$  with 2-fold cross-validation where  $\text{sign}(z)$  is defined as:

$$\text{sign}(z) = \begin{cases} 1, & z \geq 0 \\ -1, & z < 0. \end{cases}$$

- You have split the dataset  $S_{\text{training}}$  into:

$$S_1 = \{(x_1, y_1), (x_2, y_2), (x_3, y_3)\}$$

$$S_2 = \{(x_4, y_4), (x_5, y_5), (x_6, y_6)\}$$

- After training the classifier  $f(x; a, b)$  on  $S_1$ , you have obtained the parameters  $a_1 = -1, b_1 = 5$  and then try to validate the classifier on  $S_2$ .
- After training the classifier  $f(x; a, b)$  on  $S_2$ , you have obtained the parameters  $a_2 = 2, b_2 = -3$  and then try to validate the classifier on  $S_1$ .

Please finish the tasks below:

1. Calculate the **average training error** in the 2-fold cross-validation.

**Note:** The definition of average training error is the mean of the error of classifier  $f(x; a_1, b_1)$  on  $S_1$  and the error of classifier  $f(x; a_2, b_2)$  on  $S_2$ .

2. Calculate the **average validation error** (i.e. the cross-validation error) in the 2-fold cross-validation.

### 3 (12 points) Shattering

In this problem, consider a classifier  $f(\mathbf{x}; a, b) = \text{sign}(a\mathbf{x}^\top \mathbf{x} - b)$  where the feature vector  $\mathbf{x} = [x_1, x_2]^\top \in \mathbb{R}^2$  and its prediction  $f(\mathbf{x}; a, b) \in \{-1, +1\}$ . Besides,  $a, b \in \mathbb{R}$  are the model parameters and  $\text{sign}(z)$  is defined as:

$$\text{sign}(z) = \begin{cases} 1, & z \geq 0 \\ -1, & z < 0. \end{cases}$$


The classifier  $f(\mathbf{x}; a, b)$  performs a binary classification on an input feature vector  $\mathbf{x}$  under model parameters  $a, b \in \mathbb{R}$  that can be learned.

In this question, please attempt to show that if the classifier  $f(\mathbf{x}; a, b)$  can shatter two points,  $\mathbf{x}_1 = [1, 1]^\top$  and  $\mathbf{x}_2 = [-2, -3]^\top$ .


- If you think  $f(\mathbf{x}; a, b)$  **can** shatter  $\mathbf{x}_1$  and  $\mathbf{x}_2$ , you need to show that for each possible label configuration  $(y_1, y_2) \in \{(+1, +1), (-1, -1), (+1, -1), (-1, +1)\}$  of  $\mathbf{x}_1$  and  $\mathbf{x}_2$ , there exists a classifier  $f(\mathbf{x}; a, b)$  that classifies the  $\mathbf{x}_1$  and  $\mathbf{x}_2$  correctly, and you should illustrate each classifier by the following steps:
  - Draw the decision boundary of the classifier.
  - Shade the area where the classifier makes the positive prediction.

You can use the figures of the coordinate system in the next page to show your drawings.


- If you think classifier  $f(\mathbf{x}; a, b)$  **cannot** shatter  $\mathbf{x}_1$  and  $\mathbf{x}_2$ , please explain the reason.




$$(y_1, y_2) = (+1, +1)$$



$$(y_1, y_2) = (-1, -1)$$



$$(y_1, y_2) = (+1, -1)$$



$$(y_1, y_2) = (-1, +1)$$

## 4 (10 points) SVM: Gradient

Given a training dataset  $S_{\text{training}} = \{(\mathbf{x}_i, y_i)\}, i = 1, \dots, n\}$ , we wish to optimize the loss  $\mathcal{L}(\mathbf{w}, b)$  of a linear SVM classifier:

$$\mathcal{L}(\mathbf{w}, b) = \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_{i=1}^n (1 - y_i(\mathbf{w}^T \mathbf{x}_i + b))_+ \quad (1)$$

where  $(z)_+ = \max(0, z)$  is called the rectifier function and  $C$  is a scalar constant.

The optimal weight vector  $\mathbf{w}^*$  and the bias  $b^*$  used to build the SVM classifier are defined as follows:

$$\mathbf{w}^*, b^* = \arg \min_{\mathbf{w}, b} \mathcal{L}(\mathbf{w}, b) \quad (2)$$

In this problem, we attempt to obtain the optimal parameters  $\mathbf{w}^*$  and  $b^*$  by using a standard gradient descent algorithm.

**Hint:** To derive the derivative of  $\mathcal{L}(\mathbf{w}, b)$ , please consider two cases:

$$(a) \ 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 0, \quad (b) \ 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b) < 0$$

1. Derive the derivative:  $\frac{\partial \mathcal{L}(\mathbf{w}, b)}{\partial \mathbf{w}}$ .

2. Derive the derivative:  $\frac{\partial \mathcal{L}(\mathbf{w}, b)}{\partial b}$ .

## 5 (10 points) SVM: Margin

As shown in the Figure 3, we have the decision boundary (marked as a black line) defined as Eq. 3 given  $\mathbf{w}, b$ :

$$\mathbf{w}^T \mathbf{x} + b = 0 \quad (3)$$

In parallel to the decision boundary, we have the positive plane (marked as a red line) defined as Eq. 4 and the negative plane (marked as a blue line) defined as Eq. 5:

$$\mathbf{w}^T \mathbf{x} + b = +1 \quad (4)$$

$$\mathbf{w}^T \mathbf{x} + b = -1 \quad (5)$$

We pick an arbitrary point  $\mathbf{x}^-$  on the negative plane, and draw a purple line that passes  $\mathbf{x}^-$  and is perpendicular to the negative plane. The intersection between this purple line and the positive plane can be denoted as  $\mathbf{x}^+$ . Thus, we have the following Eq. 6 that indicates the relation between  $\mathbf{x}^+$  and  $\mathbf{x}^-$ :

$$\mathbf{x}^+ = \mathbf{x}^- + \lambda \mathbf{w} \quad (6)$$

where  $\lambda \in \mathbb{R}$  is an undetermined scalar. The margin  $M$  is defined as the distance between the positive and the negative planes, which can be calculated from Eq. 7:

$$M = \|\mathbf{x}^+ - \mathbf{x}^-\|_2 = \sqrt{\langle \lambda \mathbf{w}, \lambda \mathbf{w} \rangle} \quad (7)$$




Figure 3: The decision boundary, the positive plane and the negative plane.


Please derive the following according to Eq. 7:

$$M = \frac{2}{\sqrt{\langle \mathbf{w}, \mathbf{w} \rangle}}$$

**Hint:** You can firstly represent  $\lambda$  in the form of  $\mathbf{w}$  by using Eq. 4, 5, 6.

## 6 (10 points) K Nearest Neighbors

Consider a training dataset  $S_{\text{training}} = \{(\mathbf{x}_i, y_i), i = 1, 2, \dots, 8\}$  where each data point  $(\mathbf{x}, y)$  has a feature vector  $\mathbf{x} = [x_1, x_2]^\top$  and the corresponding label  $y \in \{-1, +1\}$ . The points with the corresponding labels in the dataset are shown in the figure below.



In the figure above, the index  $i$  for each training example  $\mathbf{x}_i$  is given in **bold** near the point. You are asked to predict the label of a point  $\mathbf{x}_{\text{pred}} = [1, 0]^\top$  shown in the figure as a triangle  $\blacktriangle$  using  $k$  nearest neighbors ( $k$ -NN) method under **Euclidean distance**.

1. List the indices of all the data points in  $S_{\text{training}}$  and their corresponding labels.
2. Determine the predicted label for  $\mathbf{x}_{\text{pred}}$  using the  $k$ -NN with different  $k$ .
  - (a)  $k = 1$ .
  - (b)  $k = 3$ .
  - (c)  $k = 5$ .



## 7 (20 points) Coding: K Nearest Neighbors

In this problem, you need to implement the  $k$  nearest neighbors ( $k$ -NN) algorithm and apply it to the binary classification. Here we use the modified Iris dataset  $S = \{(\mathbf{x}_i, y_i)\}$  where each feature vector  $\mathbf{x} \in \mathbb{R}^2$  and label  $y \in \{-1, +1\}$ . You are **not** allowed to use `sklearn.neighbors.KNeighborsClassifier()` in your code, but you can use it to validate your implementation.

- Load the modified Iris dataset. The dataset  $S$  is split to three subsets: The training set  $S_{\text{training}}$ , the validation set  $S_{\text{validation}}$  and the test set  $S_{\text{test}}$ . In the code, we use `X_train`, `Y_train` for the feature vectors and labels of the training set respectively. Similar notations are also used for the validation and the test sets.
- Implement  $k$ -NN algorithm in 3 steps.
  1. For each feature vector  $\mathbf{x}$  you are predicting a label, you need to calculate the distances between **this feature vector  $\mathbf{x}$**  and **all the feature vectors in the training set  $S_{\text{training}}$** .
  2. Then sort all distances in ascending order and pick the labels for the  $k$  minimum distances.
  3. Count the number of negative labels  $N_{y=-1}$ , and the number of the positive labels  $N_{y=+1}$  from  $k$  labels picked in step 2. Use the following decision rule to predict label  $\hat{y}$  for each feature vector  $\mathbf{x}$ :

$$\hat{y} = \begin{cases} +1, & N_{y=-1} < N_{y=+1}, \\ -1, & N_{y=-1} \geq N_{y=+1}. \end{cases}$$

Here we assume **Euclidean distance** as the distance metric. For more details, please refer to the code and the corresponding part in the slides.

- Use the validation set to obtain optimal  $k^*$ . In  $k$ -NN, there is a hyper-parameter  $k$  which adjusts the number of nearest neighbors. You would need to perform a grid search on the following list of  $k$ :

$$k \in \{1, 2, 3\}$$

For each  $k$ , you need to form a  $k$ -NN classifier with the training set  $S_{\text{training}}$ . Then, use the classifier to make predictions on the validation set  $S_{\text{validation}}$  and calculate the error  $e_{\text{validation}}$ . We aim to obtain the best hyper-parameter  $k^*$  corresponding to the **minimum validation error**  $e_{\text{validation}}^*$  among all  $k$ s.

- Use the obtained classifier corresponding to the best hyper-parameter  $k^*$  to calculate the test error  $e_{\text{test}}$  on test set  $S_{\text{test}}$ .

Please download the notebook `knn.ipynb` from the course website and fill in the missing blanks. Follow the instructions in the skeleton code and report:

1. Your code.
2. Plot of validation set along with decision boundary (implicitly shown in the background) corresponding to each  $k$ .
3. Validation error corresponding to each  $k$ .
4. The best hyper-parameter  $k^*$  corresponding to minimum validation error  $e_{\text{validation}}^*$ .
5. Test error  $e_{\text{test}}$  corresponding to the best hyper-parameter  $k^*$ .

## 8 (20 points) Coding: Decision Tree

In this problem, you need to implement the decision tree algorithm and apply it to the binary classification. Here we use the Ionosphere dataset  $S = \{(\mathbf{x}_i, y_i)\}$  where each feature vector  $\mathbf{x} \in \mathbb{R}^{34}$  and label  $y \in \{-1, +1\}$ . You are allowed to use the functions from `scikit-learn` in this question.

- Load the Ionosphere dataset. The dataset  $S$  is split to two subsets: The training set  $S_{\text{training}}$  and the test set  $S_{\text{test}}$ . In the code, we use `X_train`, `Y_train` for the feature vectors and labels of the training set respectively. Similar notations are also used for the test set.
- Train the decision tree classifier with the **entropy criterion**. In the decision tree, there is a hyper-parameter  $D$  which controls the maximum depth. You would need to perform a grid search on the following list of  $D$ :

$$D \in \{1, 2, 3, 4, 5\}$$

For each  $D$ , you need to form a decision tree classifier with the training set  $S_{\text{training}}$ . Specifically, you need to conduct a **10-fold** cross-validation on  $S_{\text{training}}$  and calculate the cross-validation error  $\bar{e}$  (i.e. average validation error over the splits in cross-validation). We aim to obtain the best hyper-parameter  $D^*$  corresponding to the **minimum cross-validation error**  $\bar{e}^*$  among all  $D$ s.

- Use the obtained classifier corresponding to the best hyper-parameter  $D^*$  to calculate the test error  $e_{\text{test}}$  on test set  $S_{\text{test}}$ .

Please download the notebook `decision-tree.ipynb` and the data file `ionosphere.npy` from the course website and fill in the missing blanks. Follow the instructions in the skeleton code and report:

1. Your code.
2. A heatmap of cross-validation errors corresponding to all  $D$ s.
3. The best hyper-parameter  $D^*$  corresponding to the minimum cross-validation error  $\bar{e}^*$ .
4. Test error  $e_{\text{test}}$  corresponding to the best hyper-parameter  $D^*$ .